

Convolutional Neural Nets (Convnets)

Excellent reference material:

<http://cs231n.stanford.edu/> - Stanford's Convnet class. Video of lectures available at
<https://www.youtube.com/playlist?list=PL3FW7Lu3i5JvHM8ljYj-zLfQRF3EO8sYv>

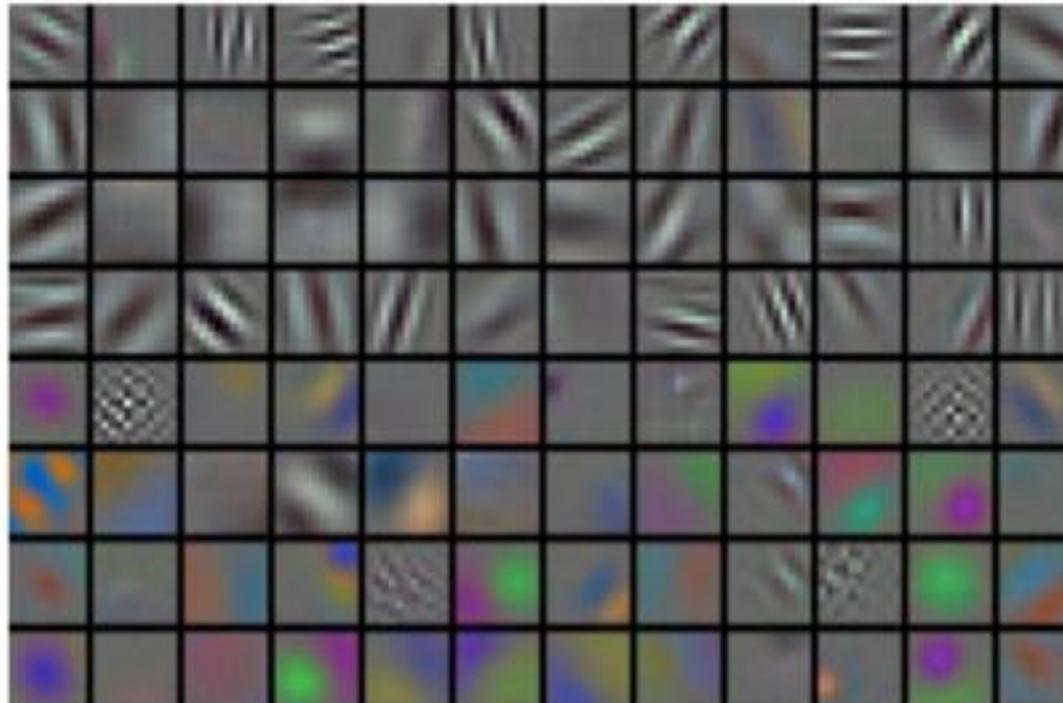
Motivation

- So far we've seen several “hand-crafted” features for describing local image structure.
- Examples
 - Derivative of Gaussian edges
 - 2nd Derivative of Gaussian edges
 - Laplacian of Gaussian edges and blobs
 - Harris Corner R score edges, corners, uniform regions

Motivation

Given a set of images with labels about what objects they contain, can we automatically learn features that are good for classifying those different objects?

Yes!



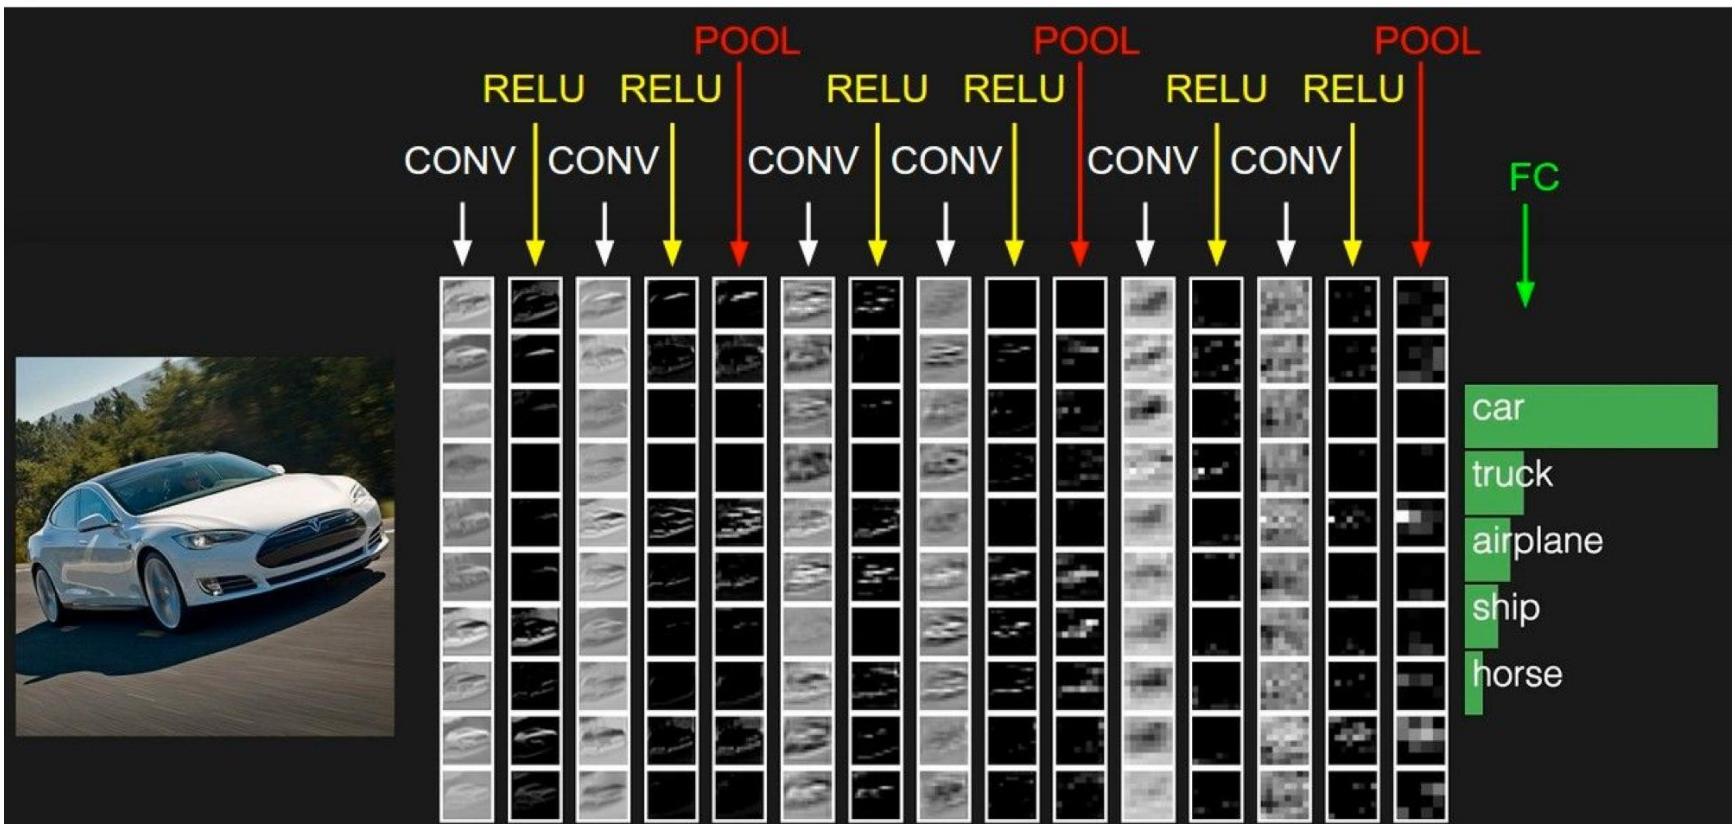
Some convolution filters learned by a neural network.

Convnets

- A very modular kind of classifier. Number of layers chained together are considered to be the “depth”. Result is a deep classifier, hence deep learning.



preview:



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 7 - 22

27 Jan 2016

Demo running in browser, at <http://cs231n.stanford.edu/>

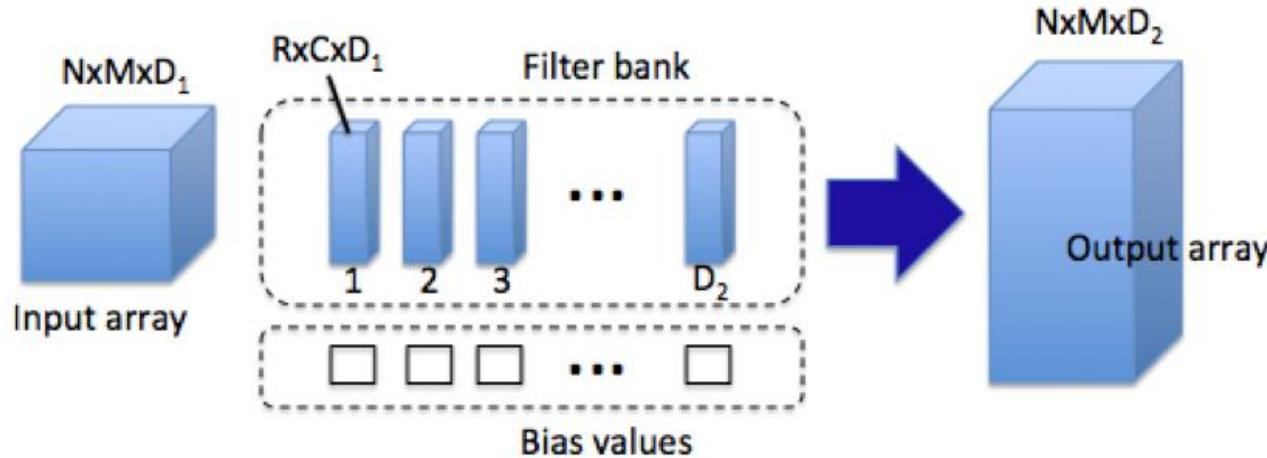
Convnet Overview

- Forming complex classifiers as long chains of simple computational building blocks.
- Building Blocks
 - Convolution (actually cross-correlation)
 - Nonlinear Activation (e.g. ReLU)
 - Pooling (e.g. max-pool)
 - Fully Connected Layers
 - Softmax
- Goal is to understand forward pass computation.
- Understanding backpropagation (how learning takes place) is outside the scope of this lecture.

Overview

- Building Blocks
 - Convolution (actually cross-correlation)
 - Nonlinear Activation (e.g. ReLU)
 - Pooling (e.g. max-pool)
 - Fully Connected Layers
 - Softmax
- Goal is to understand forward pass computation

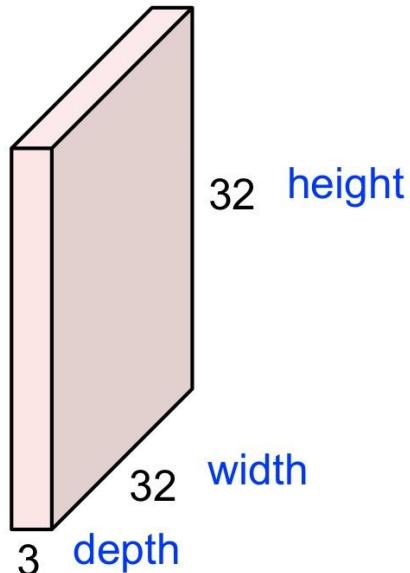
Convolution Layer



A convolution layer transforms an $N \times M \times D_1$ input array into an $N \times M \times D_2$ output array by convolving with a filter bank of D_2 linear filters and adding bias values.

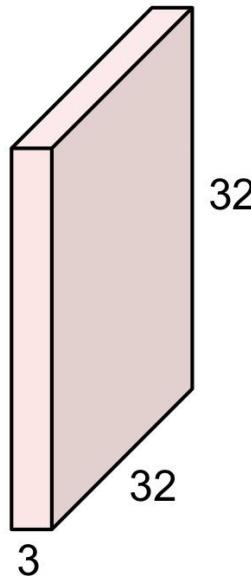
Convolution Layer

32x32x3 image



Convolution Layer

32x32x3 image



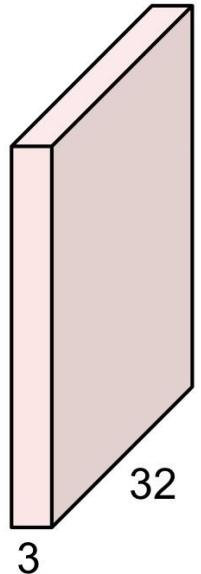
5x5x3 filter



Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer

32x32x3 image



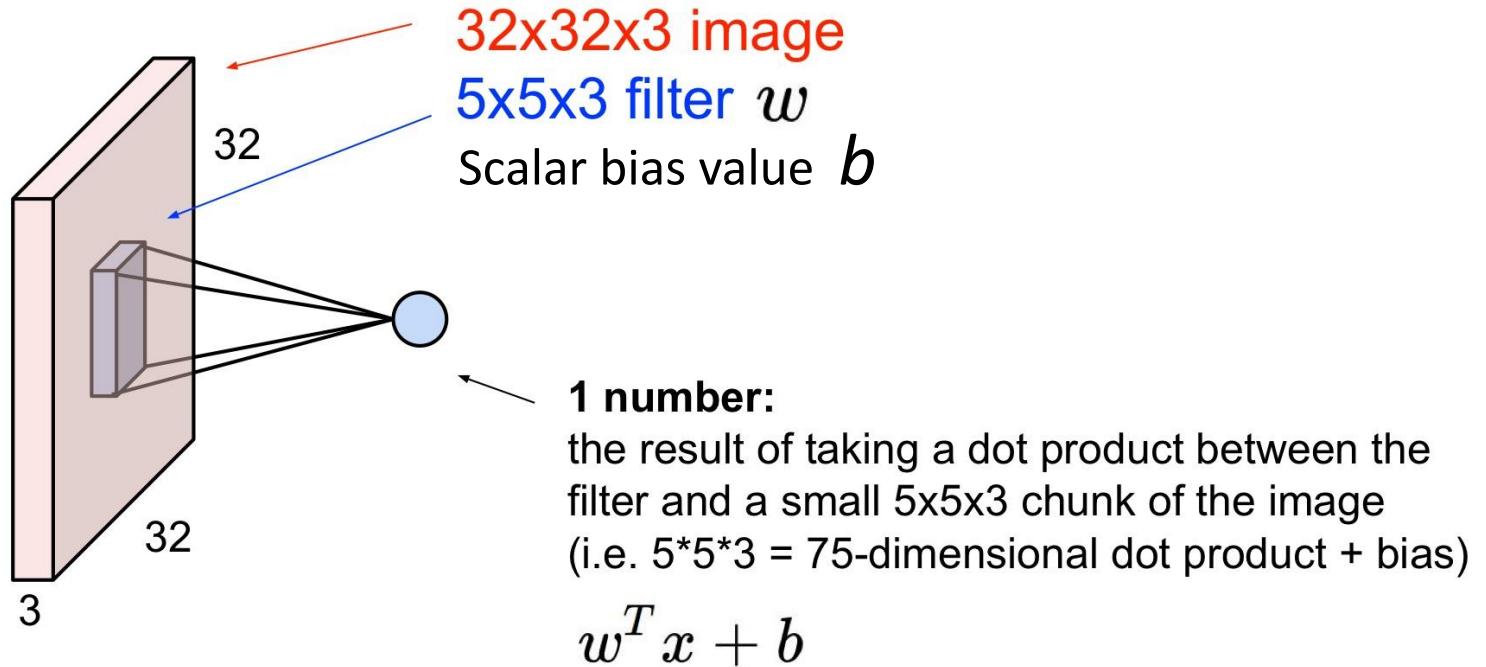
Filters always extend the full depth of the input volume

5x5x3 filter

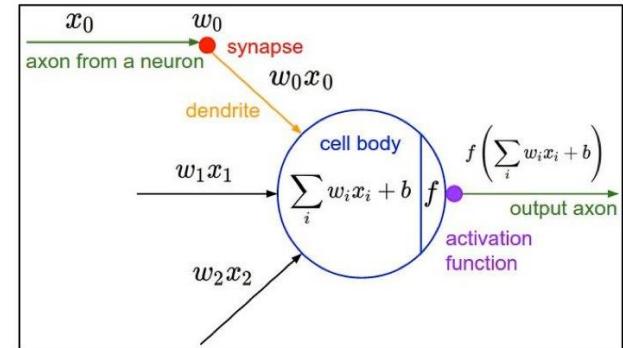
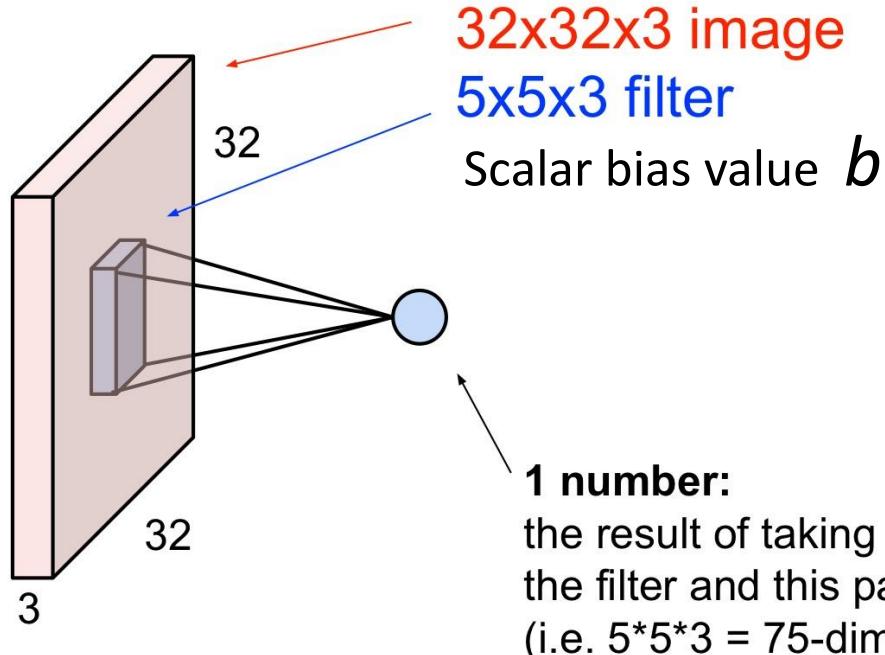


Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer



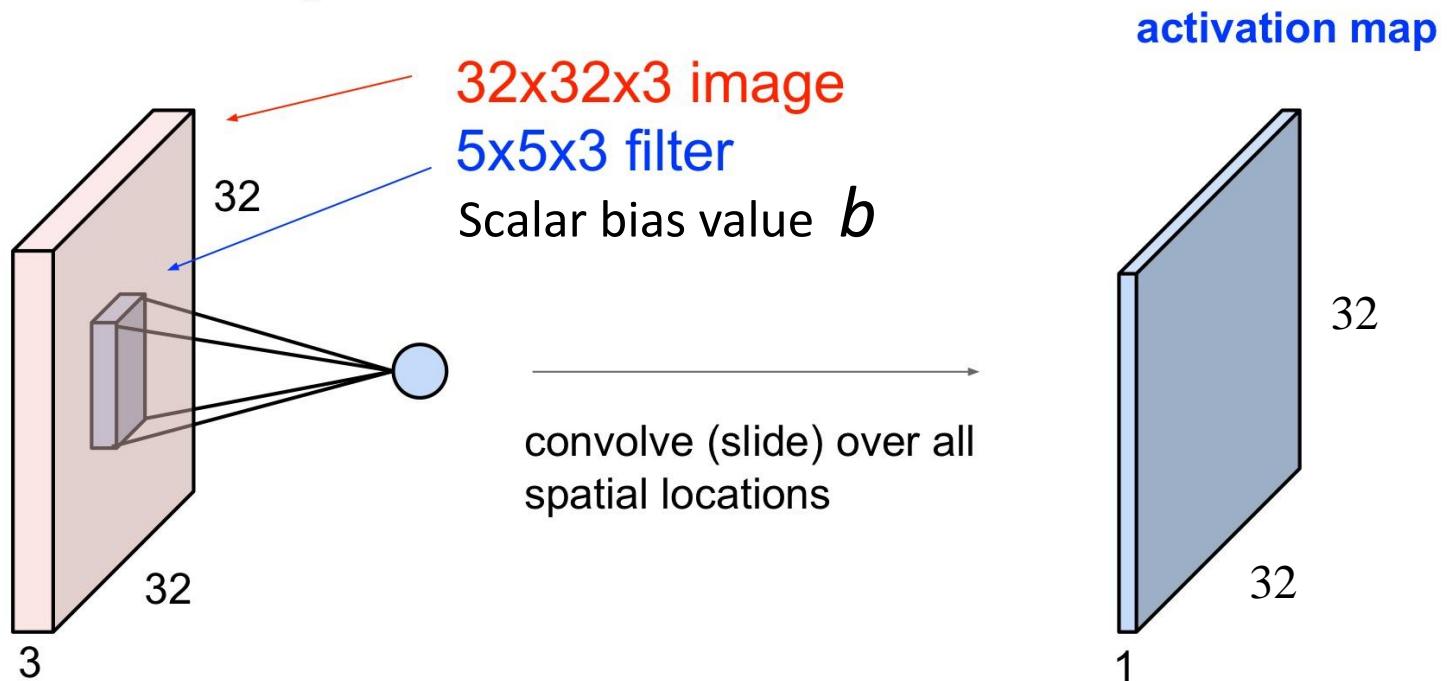
The brain/neuron view of CONV Layer



It's just a neuron with local connectivity...

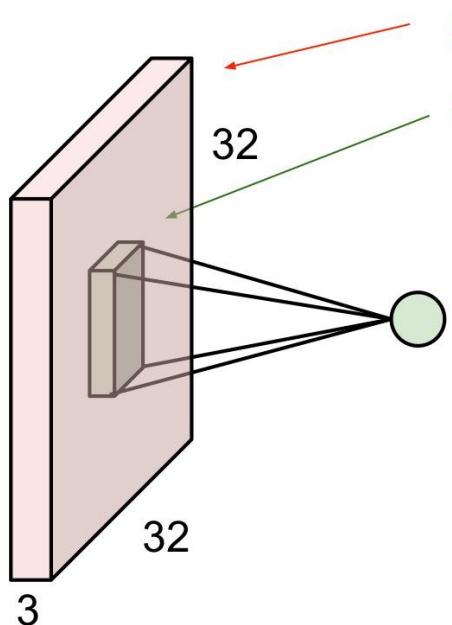
Activation function here is just identity function

Convolution Layer



Convolution Layer

Consider a second, different filter



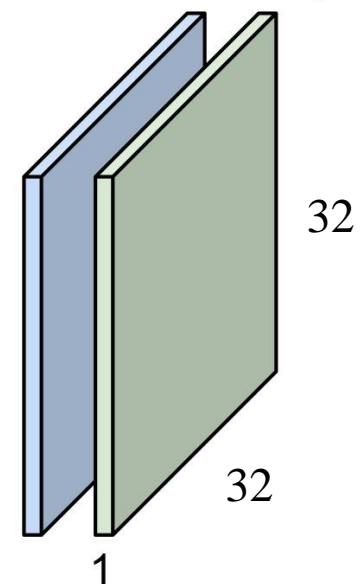
32x32x3 image

5x5x3 filter

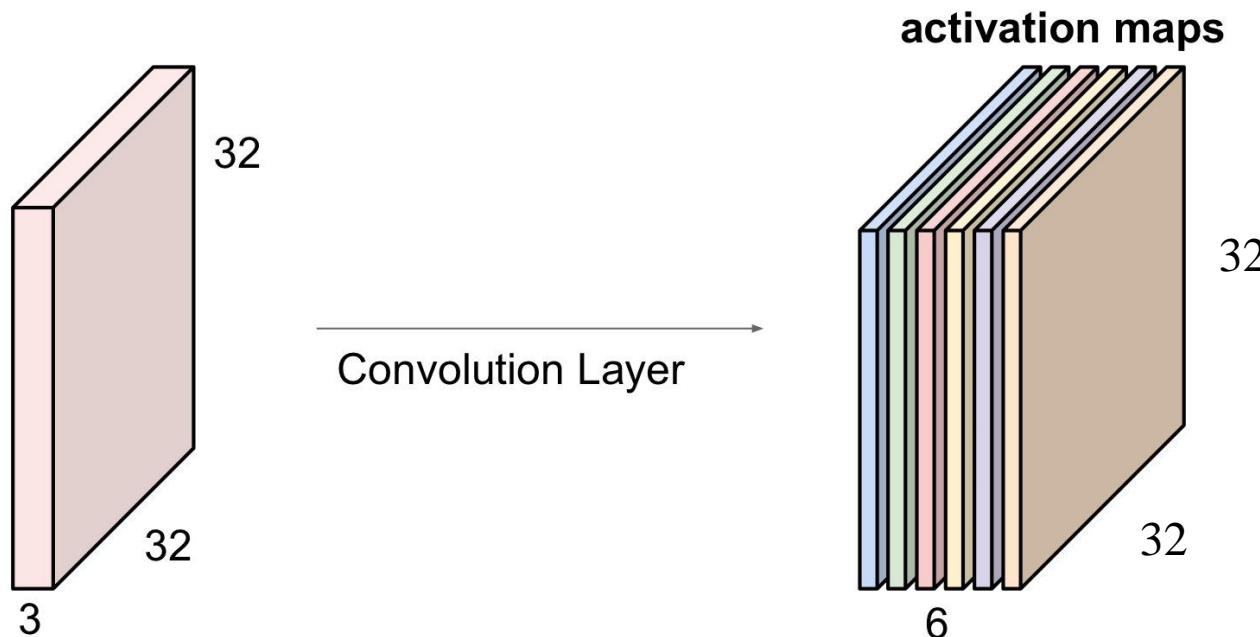
Scalar bias value b_2

convolve (slide) over all
spatial locations

activation maps



and 6 scalar bias values,
For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size $32 \times 32 \times 6$

Equivalent to the idea of applying a convolution filter bank
in order to generate a vector of feature values.

Overview

- Building Blocks
 - Convolution (actually cross-correlation)
 - Nonlinear Activation (e.g. ReLU)
 - Pooling (e.g. max-pool)
 - Fully Connected Layers
 - Softmax
- Goal is to understand forward pass computation

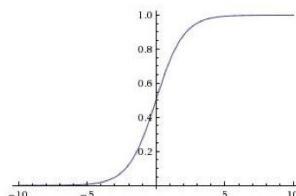
Motivation for Activation Functions

- Convolution is a linear operator
- Cascades of convolutions are also linear
- The complexity of a linear classifier's decision boundary is limited.
- Activation functions introduce nonlinearity into the system, so that it can have nonlinear decision boundaries!

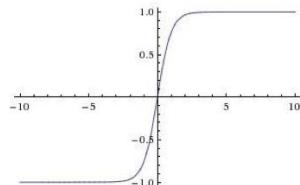
Activation Functions

Sigmoid

$$\sigma(x) = 1/(1 + e^{-x})$$

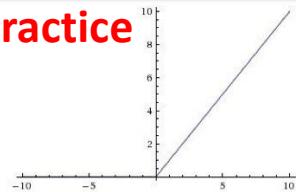


tanh tanh(x)



Most common in current practice

ReLU max(0,x)

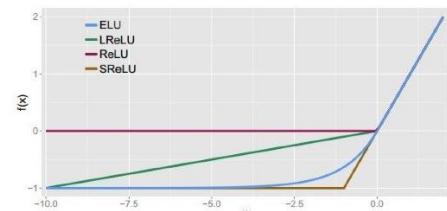


Fei-Fei Li & Andrej Karpathy & Justin Johnson

Maxout ELU

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$



Lecture 5 - 29

20 Jan 2016

ReLU clamps negative values in an array to 0. Nonnegative values remain the same.

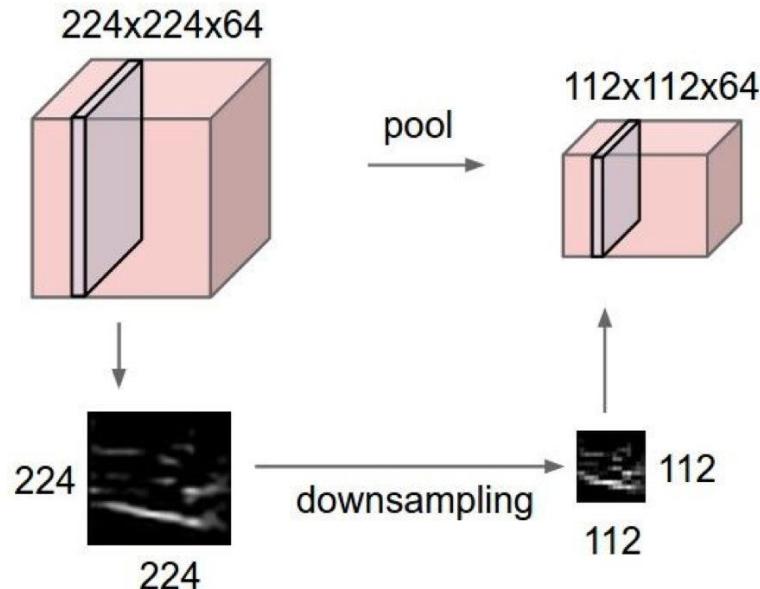
$$\begin{matrix} 5 & -1 \\ -2 & 4 \end{matrix} \longrightarrow \begin{matrix} 5 & 0 \\ 0 & 4 \end{matrix}$$

Overview

- Building Blocks
 - Convolution (actually cross-correlation)
 - Nonlinear Activation (e.g. ReLU)
 - Pooling (e.g. max-pool)
 - Fully Connected Layers
 - Softmax
- Goal is to understand forward pass computation

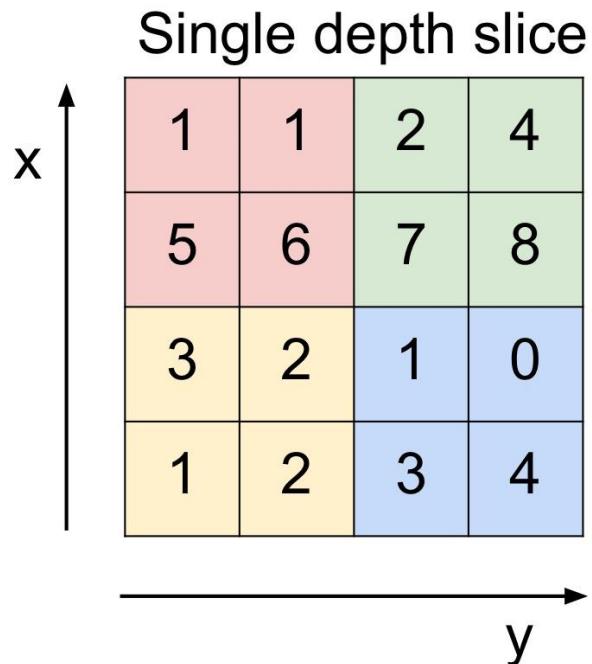
Pooling layer

- makes the representations smaller and more manageable
- operates over each activation map independently:



- Also adds resilience to small geometric deformations.
- A way to do spatial hierarchical processing (fine to coarse).

MAX POOLING

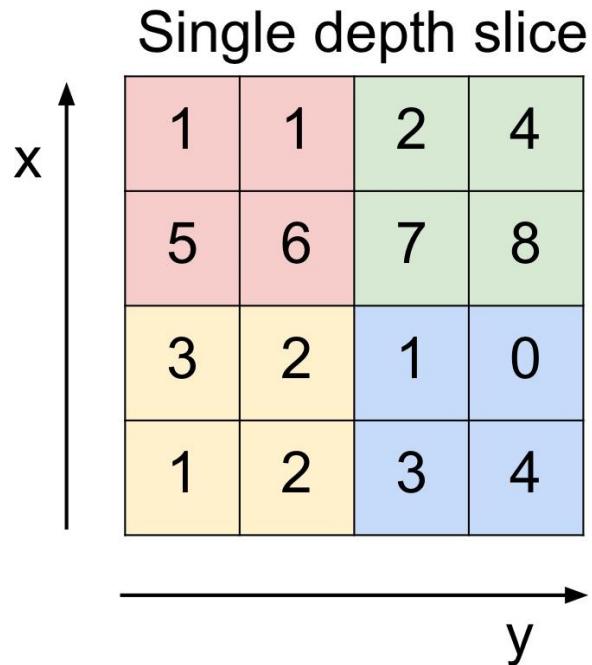


max pool with 2x2 filters
and stride 2

A 2x2 output matrix resulting from max pooling with 2x2 filters and stride 2. The matrix has two rows and two columns. The top-left cell contains the value 6 (pink), the top-right cell contains 8 (light green), the bottom-left cell contains 3 (yellow), and the bottom-right cell contains 4 (light blue).

6	8
3	4

Mean Pooling



mean
~~max~~ pool with 2x2 filters
and stride 2

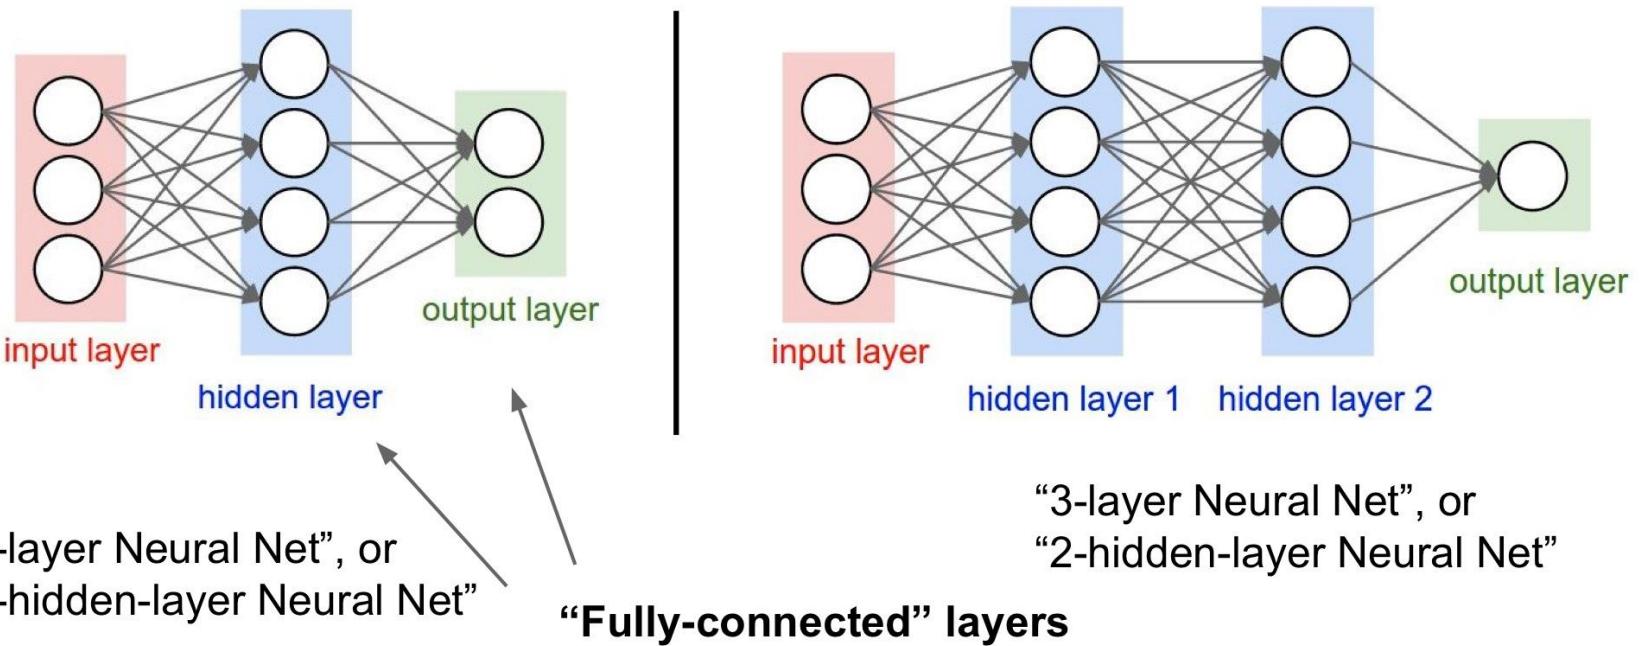
The figure shows a 2x2 output tensor with values arranged in two rows and two columns. The x-axis is labeled 'x' and points upwards, while the y-axis is labeled 'y' and points to the right. The values in the tensor are:

13/4	21/4
8/4	8/4

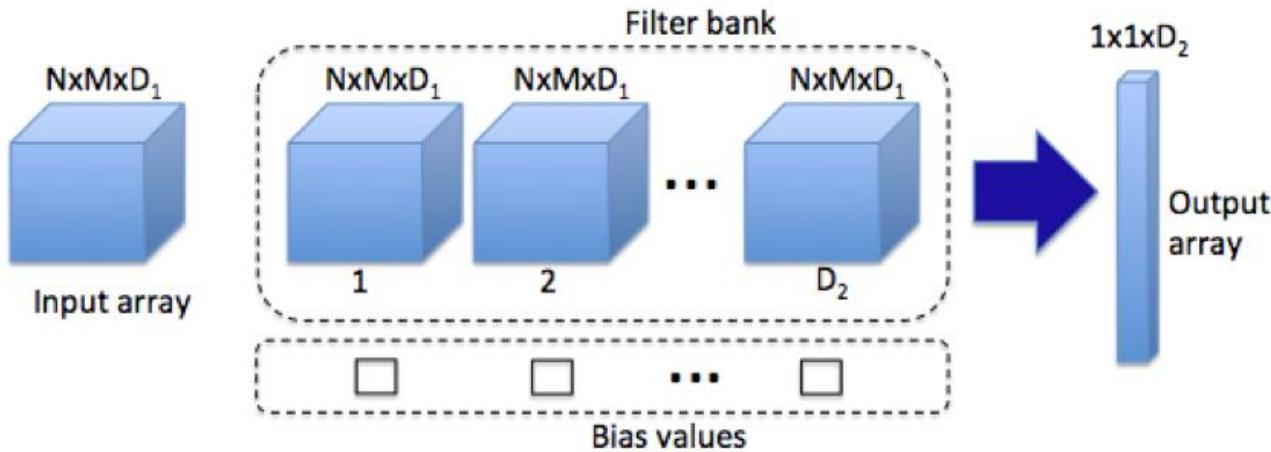
Overview

- Building Blocks
 - Convolution (actually cross-correlation)
 - Nonlinear Activation (e.g. ReLU)
 - Pooling (e.g. max-pool)
 - Fully Connected Layers
 - Softmax
- Goal is to understand forward pass computation

Neural Networks: Architectures

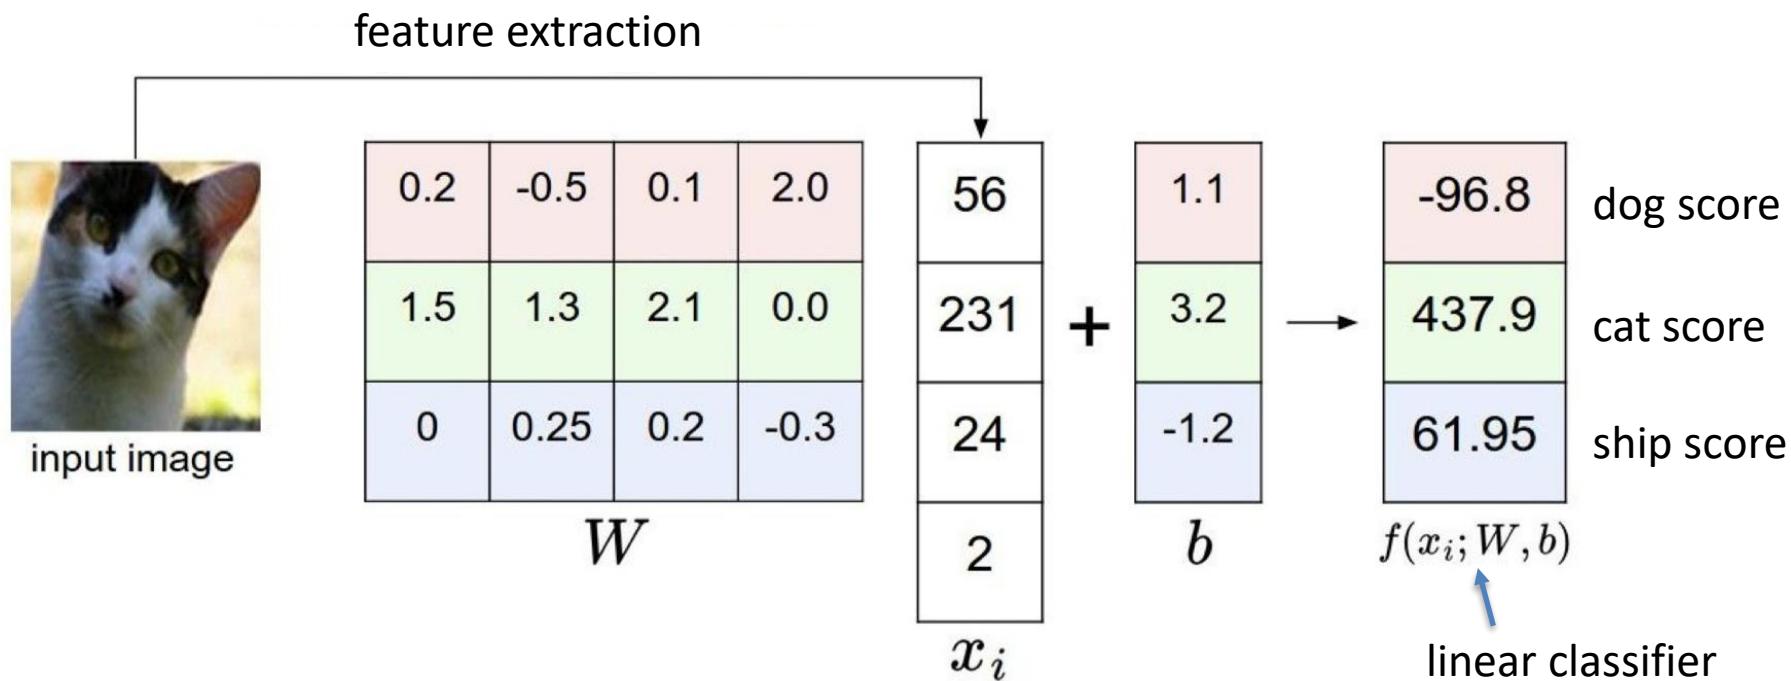


Fully Connected Layer



I like to think of fully connected layers as a special case of convolution layers where the filter is the same size as the image, and applied in only one location (no sliding).

Classifying an image using four features into three classes...



Another way of visualizing fully connected layers.
Tie in with linear classifiers!!

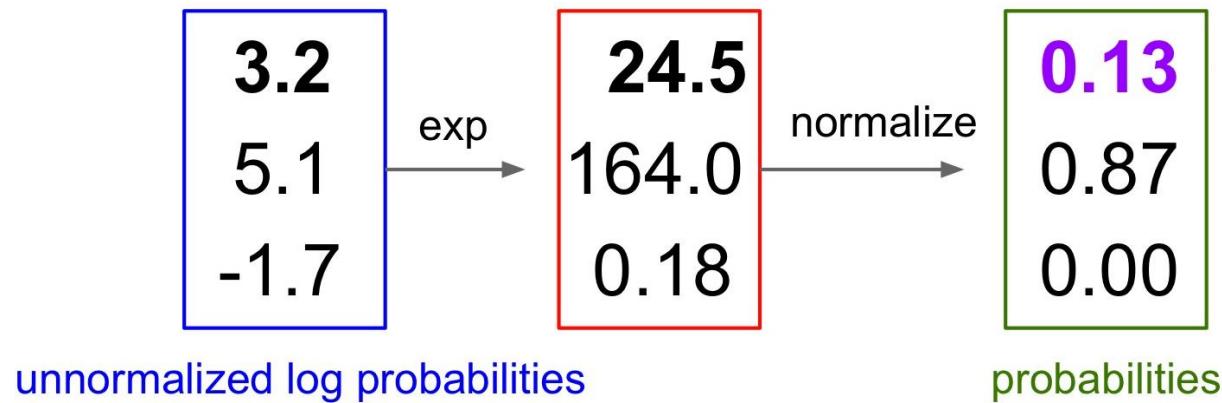
Overview

- Building Blocks
 - Convolution (actually cross-correlation)
 - Nonlinear Activation (e.g. ReLU)
 - Pooling (e.g. max-pool)
 - Fully Connected Layers
 - Softmax
- Goal is to understand forward pass computation

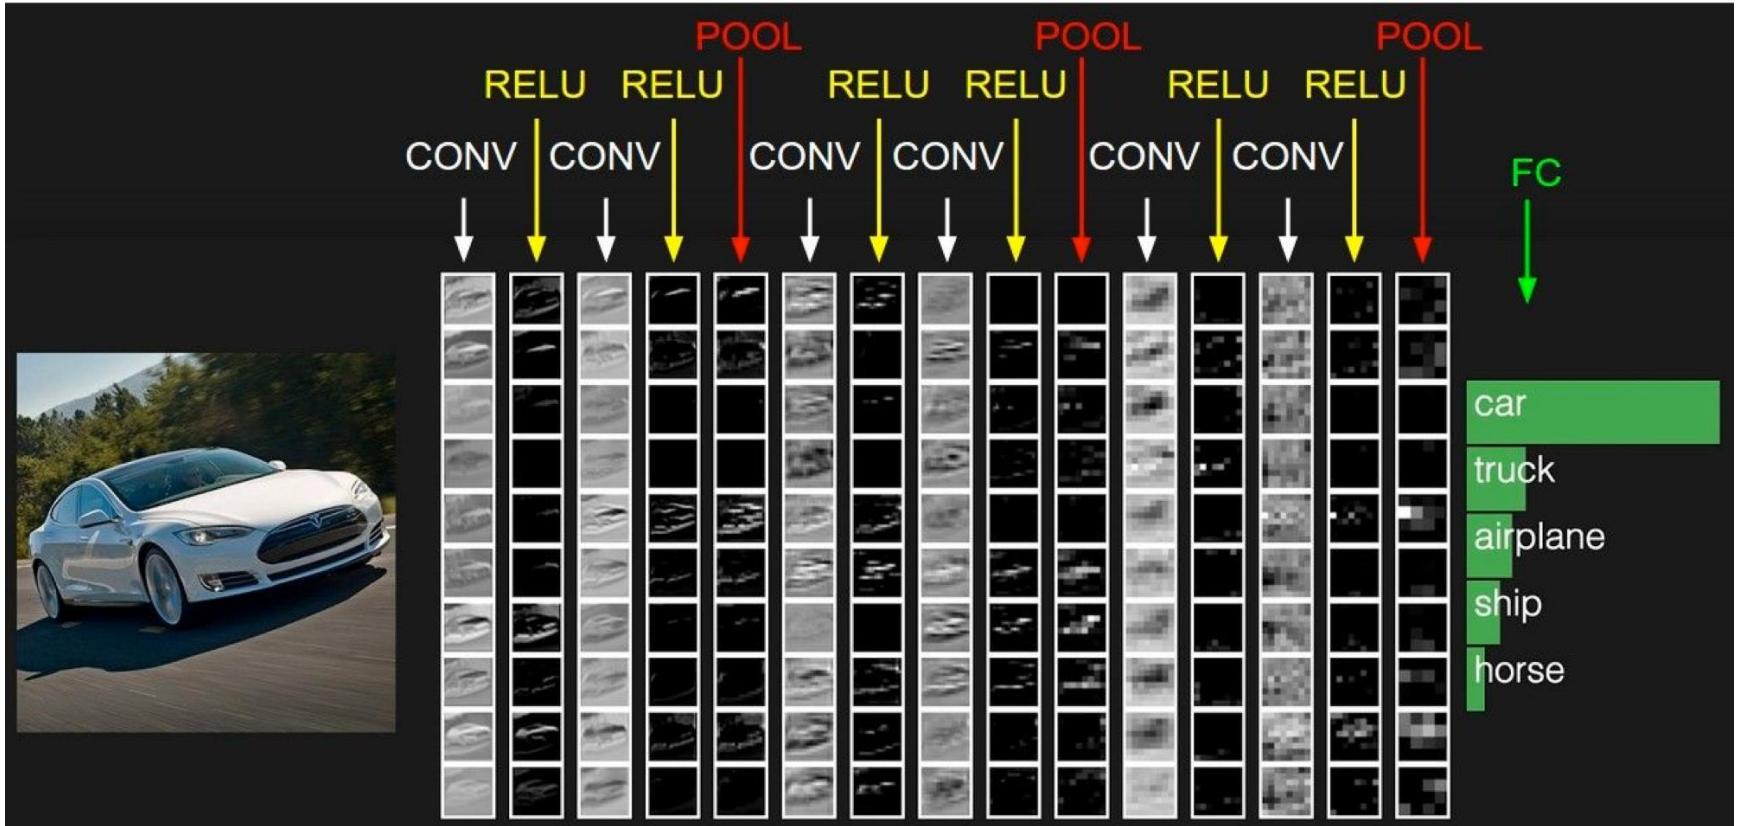
Softmax

- Turns an arbitrary set of numbers into “probabilities”.
- $p_i = \exp(x_i) / \sum_{j=1 \text{ to } n} \exp(x_j)$
- For numerical stability, it is a good idea to first subtract max value from each x_i . i.e. $x'_i = x_i - \max(\{x_j, j=1, \dots, n\})$.

Softmax



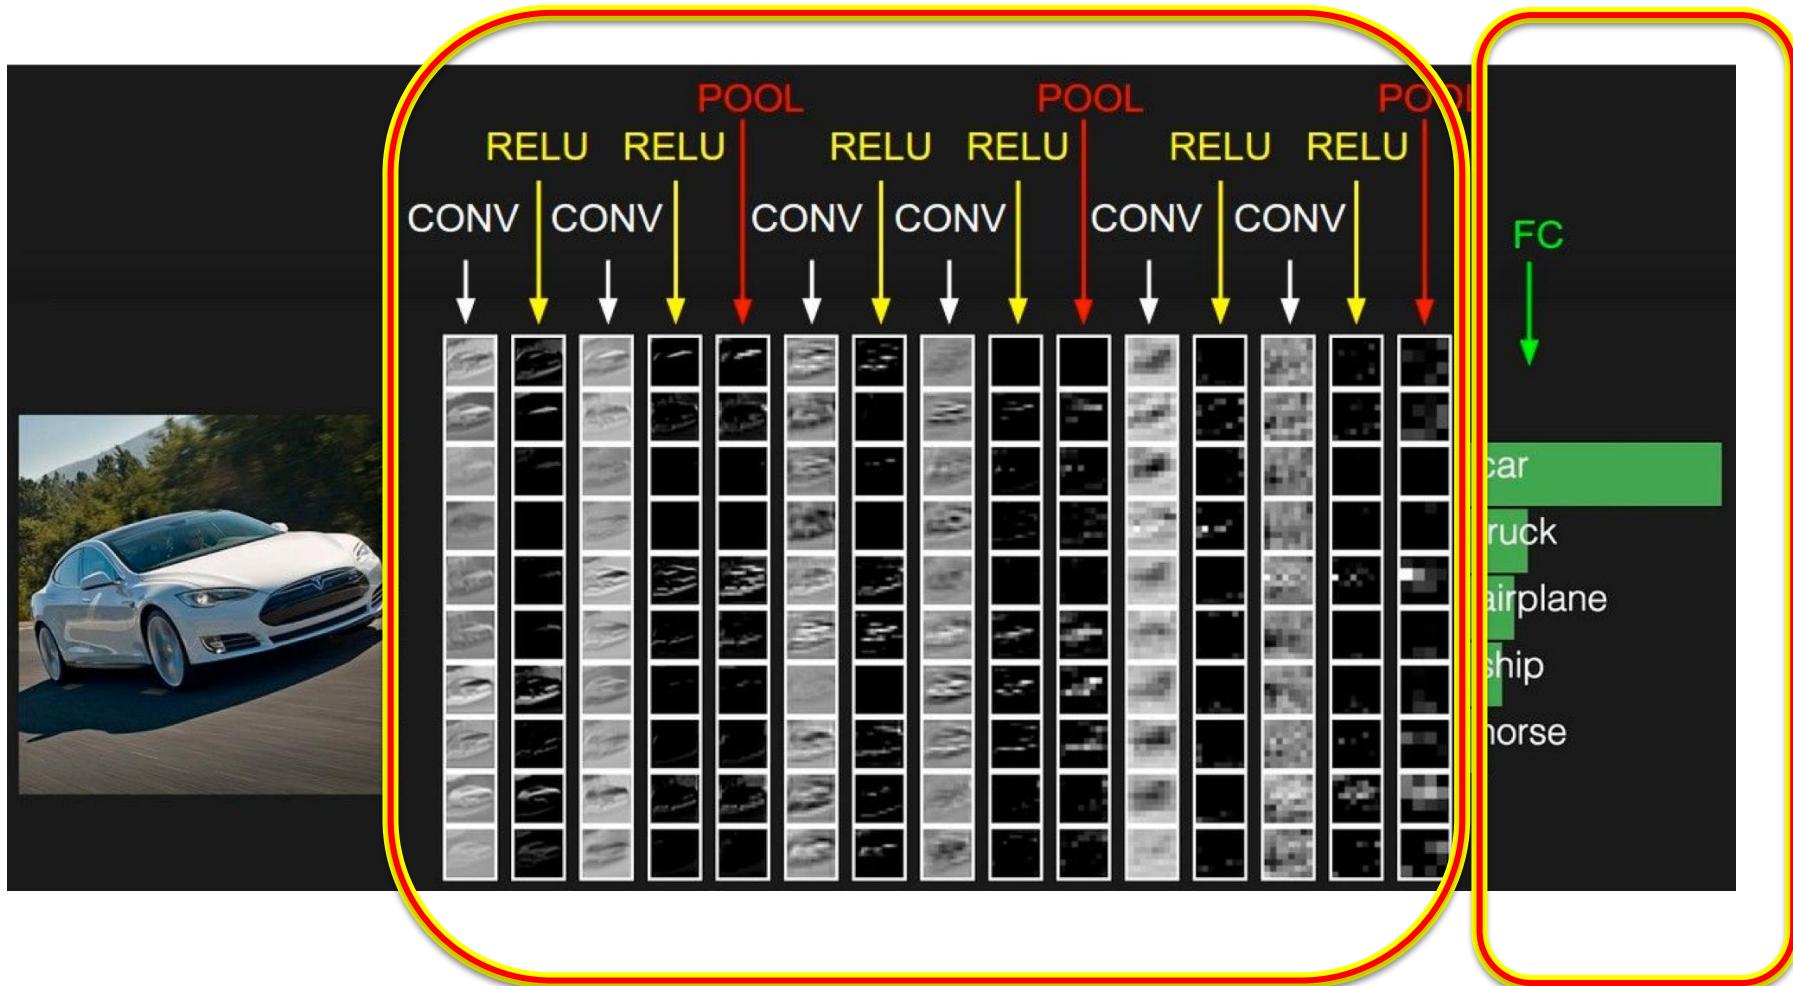
Putting it All Together



One Interpretation

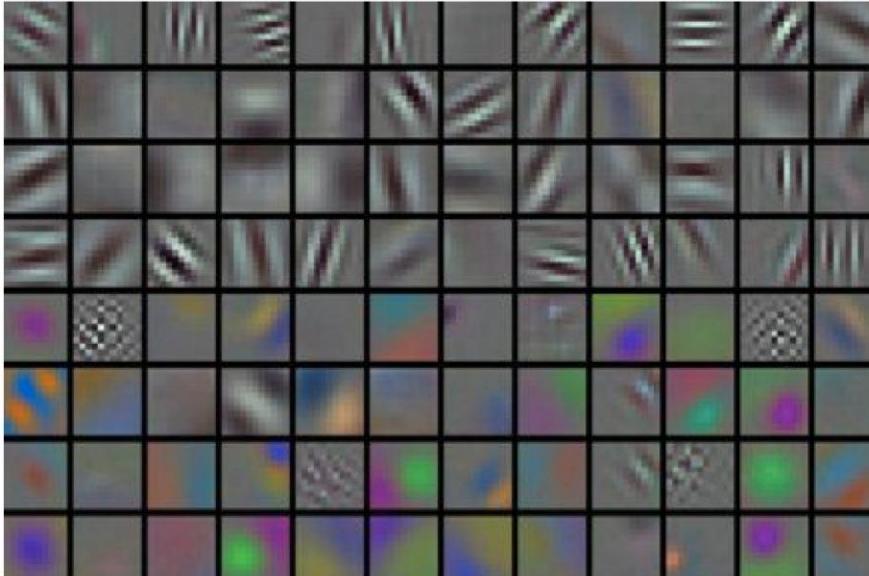
Extract features (from feature space
the linear classifier operates on)

Apply multi-class
linear classifier!

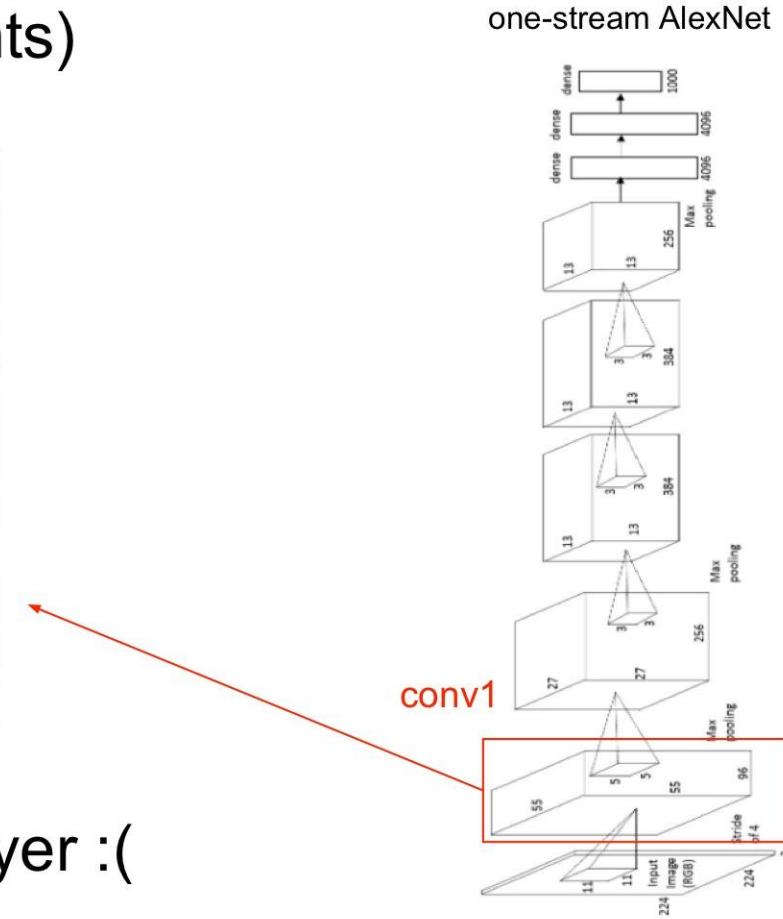


Learned Features

Visualize the filters/kernels (raw weights)



only interpretable on the first layer :(



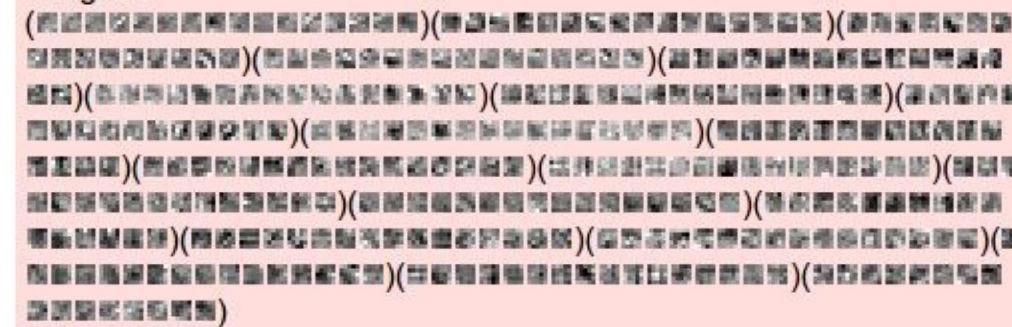
Visualize the filters/kernels (raw weights)

you can still do it for higher layers, it's just not that interesting

(these are taken from ConvNetJS CIFAR-10 demo)

Weights:


layer 1 weights

Weights:


layer 2 weights

Weights:

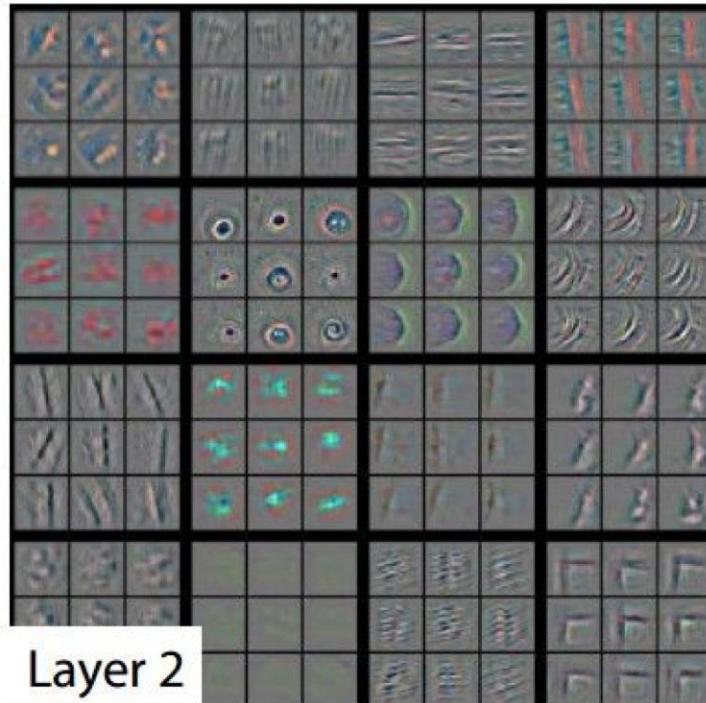

layer 3 weights

Weights in higher-level filters are rewarding or inhibiting lower-level features, not raw pixels directly.

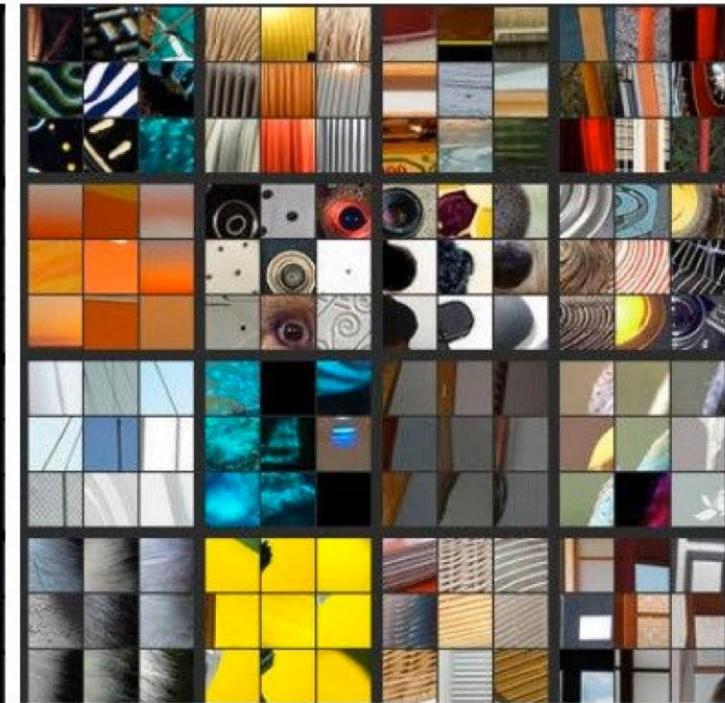
Visualizing arbitrary neurons



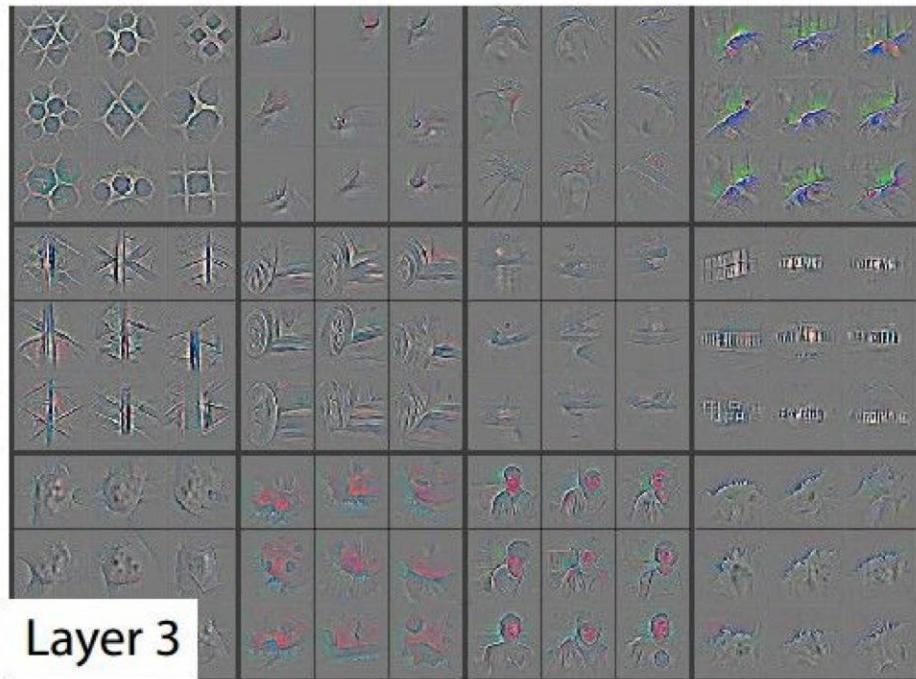
Layer 1



Layer 2



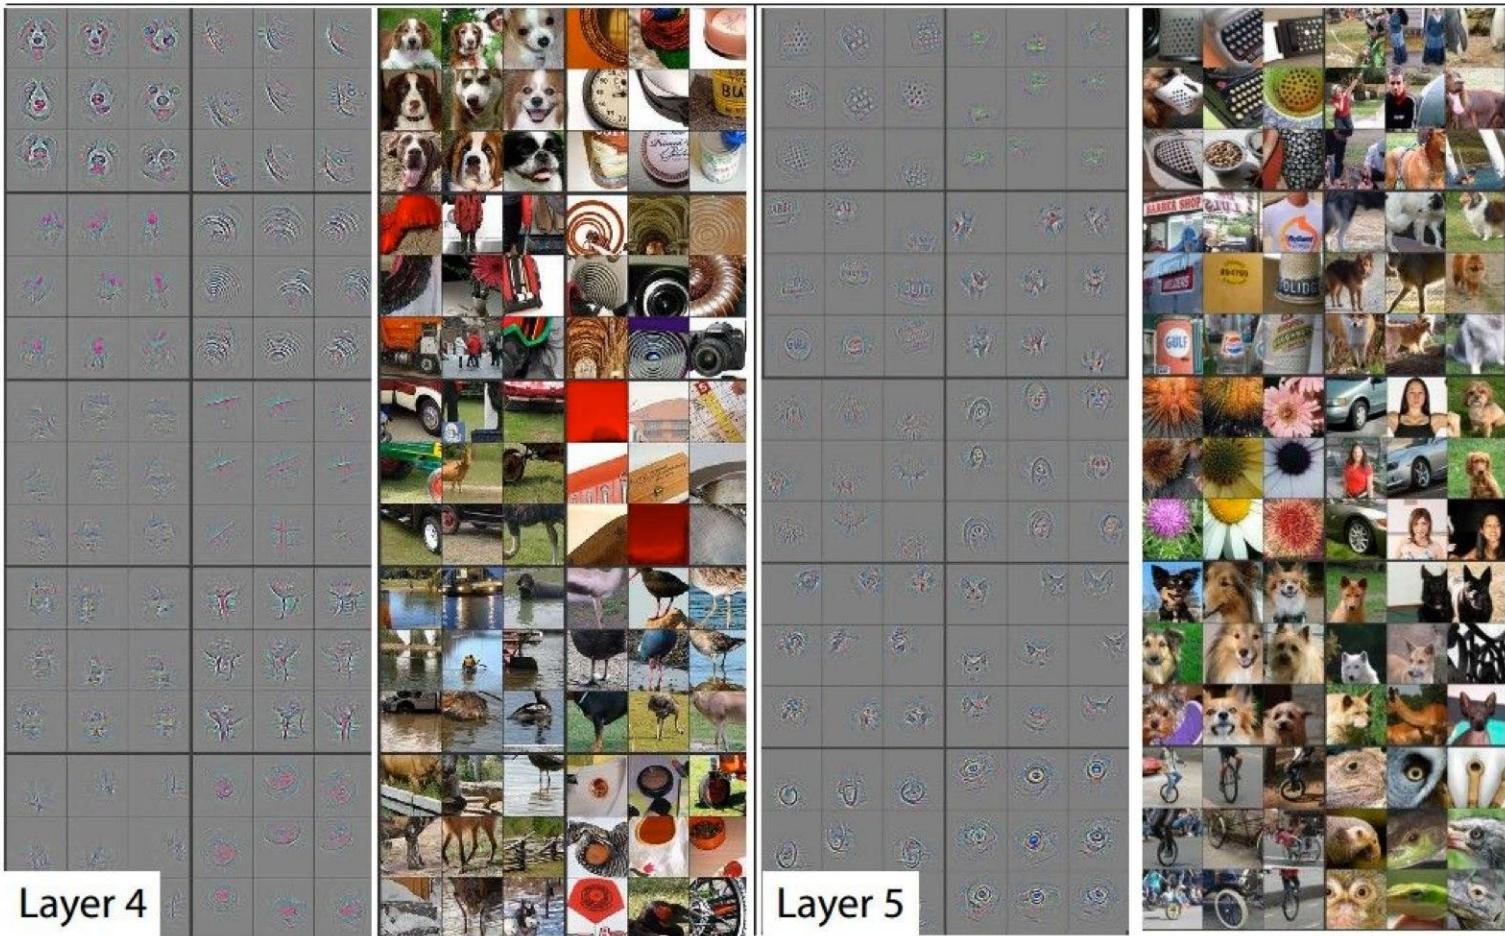
Visualizing arbitrary neurons



Layer 3



Visualizing arbitrary neurons



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Lecture 9 - 28

3 Feb 2016